

EECS 510: Social Media Mining
Spring 2016

Introduction to Deep Learning: Software Packages

Rosanne Liu
rosanne.liu@northwestern.edu
May 5, 2016

Outline

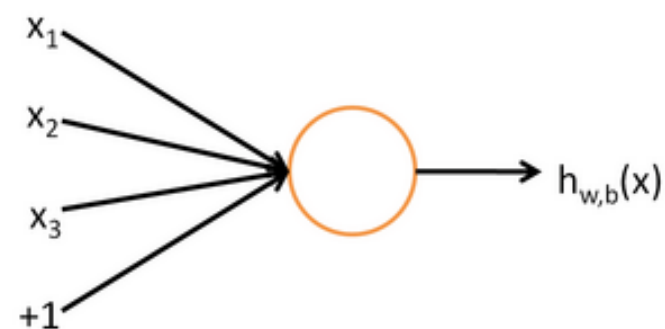
- Review
 - Deep neural networks
 - Backpropagation
- A quick start with Torch
- A quick start with Theano
- A quick start with TensorFlow

Deep learning review

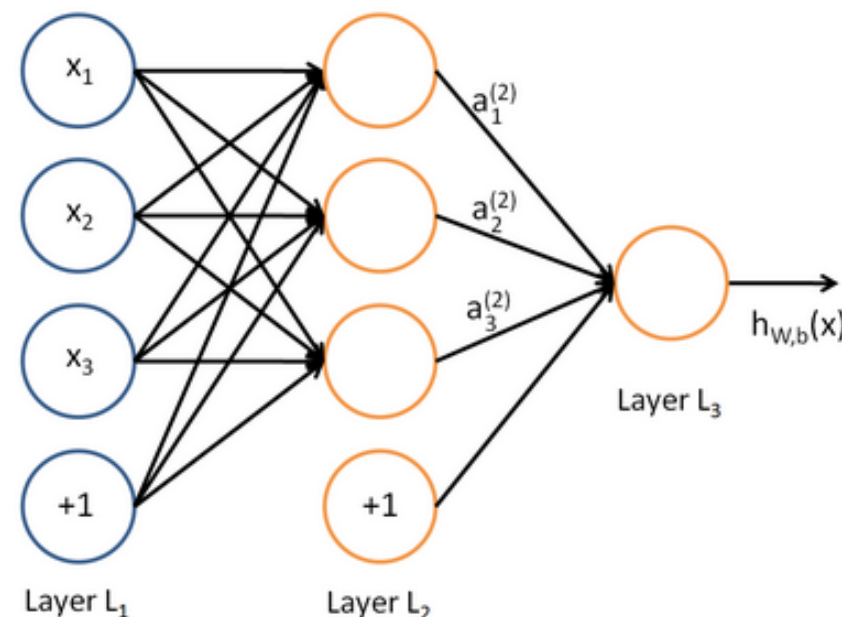
(Artificial) neural networks and training

- Artificial neural networks
 - biology inspired, supervised, weights, activation, backpropagation

A “neuron”



A small neural net



Training of a NN

Loop until tired:

1. **Sample** a batch of data.
2. **Forward** it through the network to get predictions.
3. **Backprop** the errors.
4. **Update** the weights.

Backprop

- Backprop: a way of computing *gradients* of expressions through recursive application of **chain rule**
- Problem: given a function $f(x)$ where x is a vector of inputs, we are interested in computing the gradient of f at x , i.e., $\nabla f(x)$
- Autodifferentiation
 - <http://arxiv.org/pdf/1502.05767v2.pdf>

Deep learning package zoo

- Torch
- Caffe
- Theano (keras, lasagne)
- CuDNN
- Tensorflow
- Mxnet
- Etc.

theano



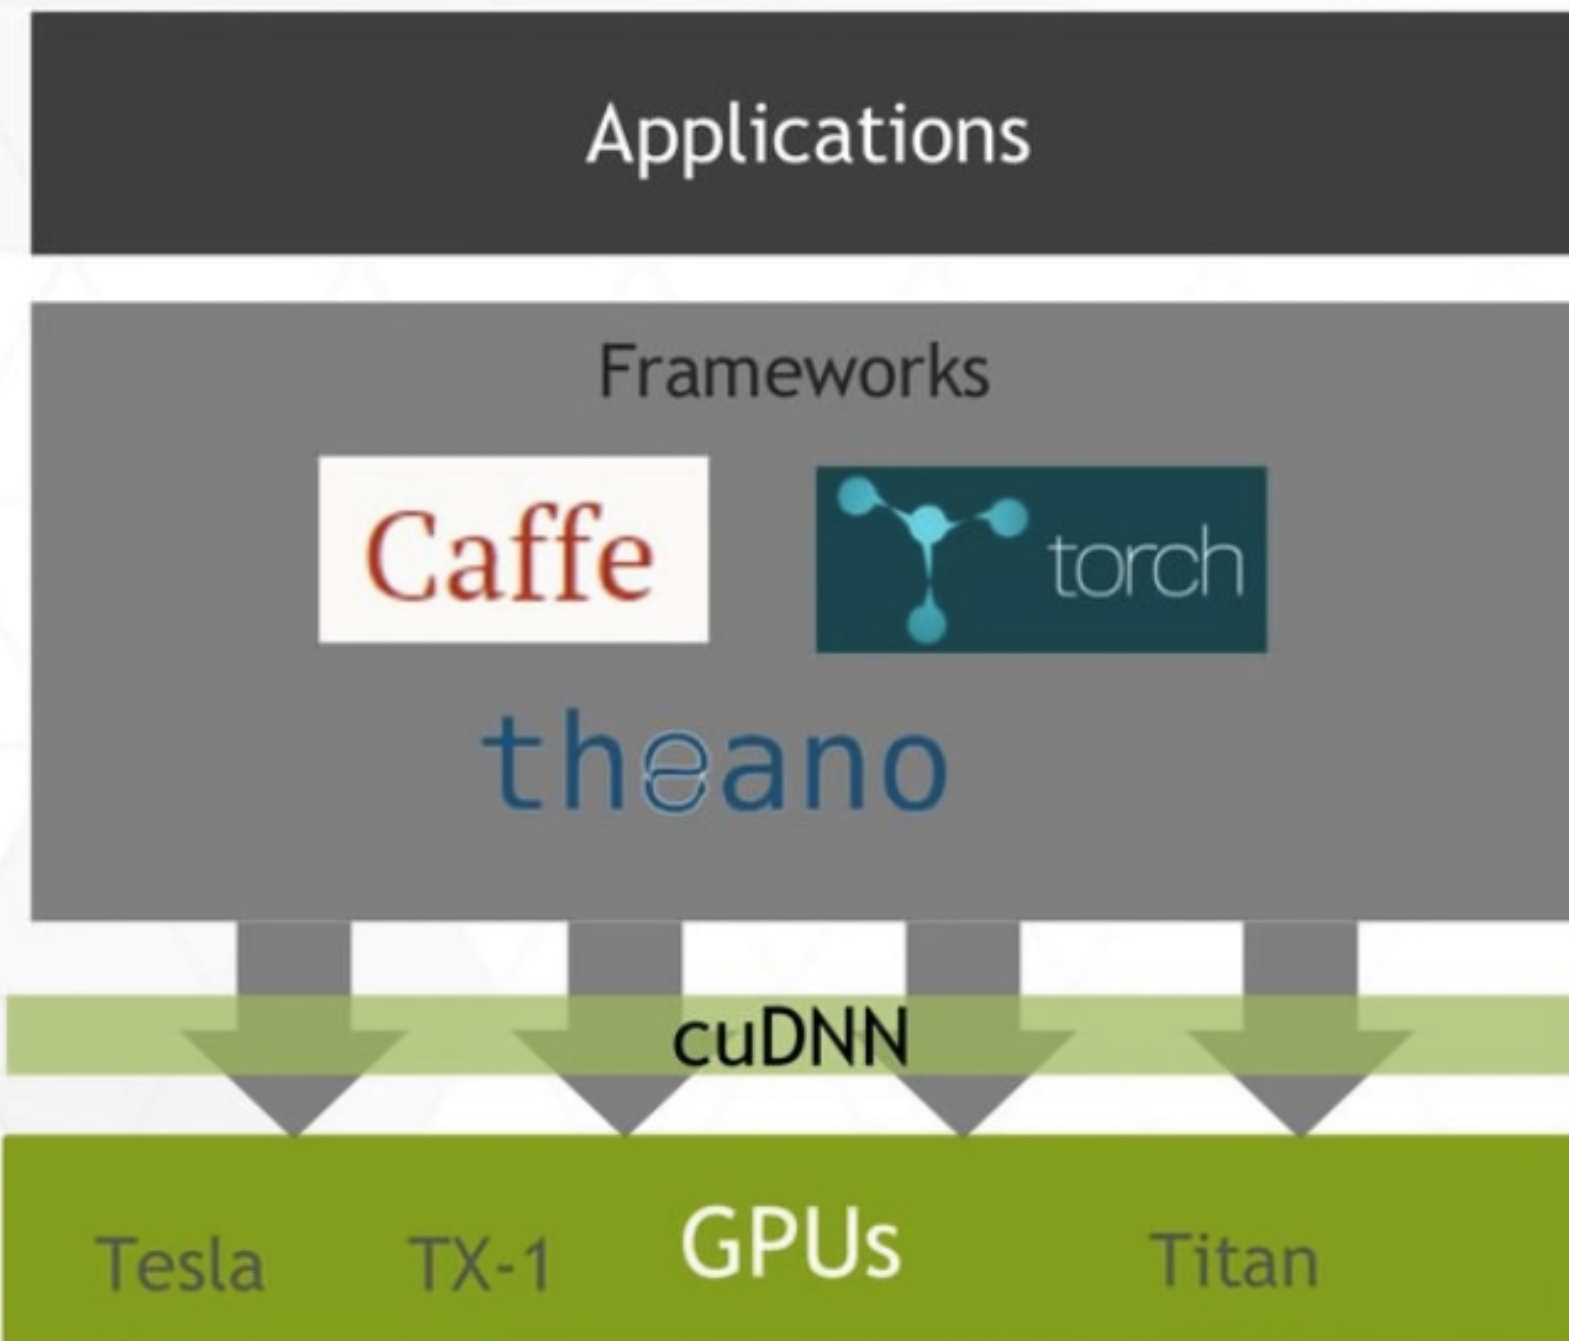
Details of package

- Backend
 - Krizhevsky (AlexNet) first built CNN in CUDA (cunn)
 - Then Nvidia took the charge and built [cudnn](#)
- Berkeley's Caffe
 - modular CNN package in C++
 - with both CPU/GPU training
- Bengio's Theano
 - Python project (w/ numpy & scipy)
 - works with GPU through cunn, cudnn
 - compile to C on-the-fly
- Facebook/NYU's Torch7
 - LuaJIT interface to C
- Google's TensorFlow
 - C with python interface

What they all have in common

- Tensor
- Symbolic differentiation
- GPU support (through backend like CuDNN)
- Open source

CuDNN is behind every package



CuDNN

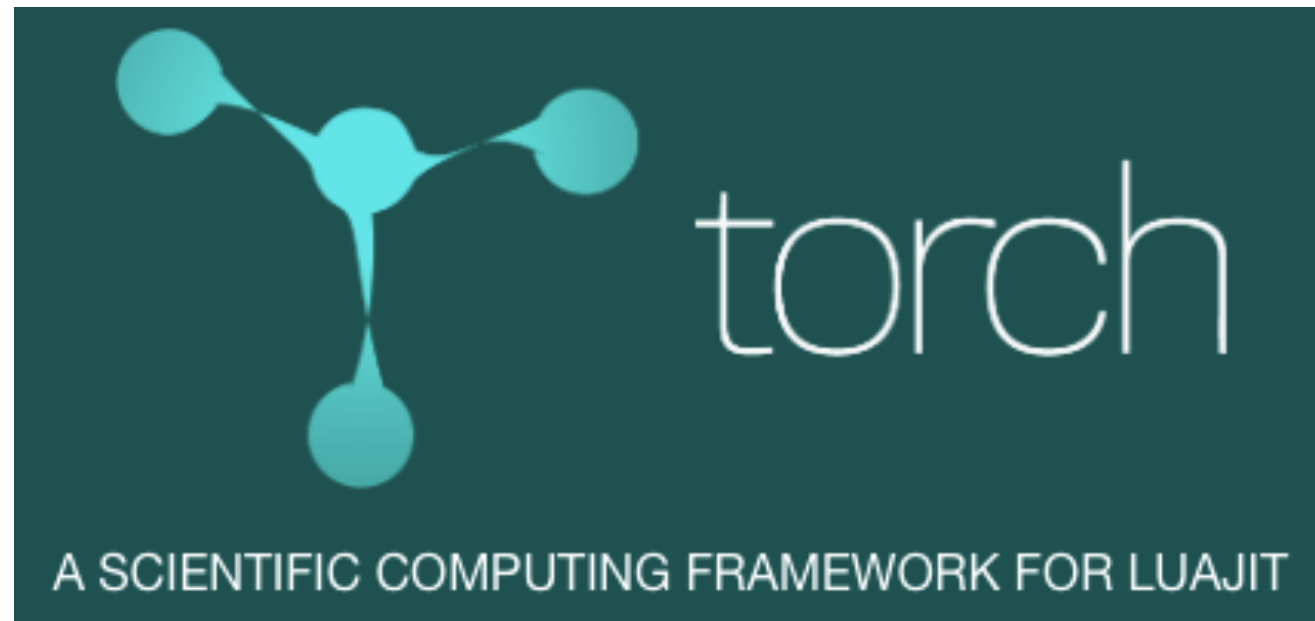
- GPU-accelerated Deep Learning subroutines
- High performance neural network training
- Accelerates major deep learning frameworks: Caffe, Theano, Torch, TensorFlow
- Up to 3.5x faster AlexNet training in Caffe than baseline GPU

Deep learning package design choices

- Model specification
 - configuration file (e.g. Caffe, CNTK)
 - Programmatic generation (e.g. Torch, Theano, TensorFlow)
- Programming language
 - Everything links to C at the bottom
 - Lua (Torch)
 - Python (Theano, TensorFlow)

A quick start with Torch

A quick start with Torch



- Project started ~2000 at Facebook, now Torch7 - 4th generation (using odd numbers only 1,3,5,7)
- Web-scale learning in speech, image and video applications
- Maintained by top researchers (@Facebook, twitter, DeepMind)
- Up until 6 days ago DeepMind has been using it



Cheatsheet

- Cheatsheet
 - <https://github.com/torch/torch7/wiki/Cheatsheet>
- Github
 - <https://github.com/torch/torch7>
- Google Group for new users and installation questions
 - <https://groups.google.com/forum/embed/?place=forum%2Ftorch7#!forum/torch7>
- **Advanced only**
 - <https://gitter.im/torch/torch7>

Tensors

- The Tensor class is the most important class in Torch (and others)
- A Tensor is a serializable, potentially multi-dimensional matrix
- The number of dimensions in Torch is unlimited (created using LongStorage)
- A scalar is a tensor $(f : \mathbb{R} \rightarrow \mathbb{R})$
- A vector is a tensor $(f : \mathbb{R}^n \rightarrow \mathbb{R})$
- A matrix is a tensor $(f : \mathbb{R}^n \times \mathbb{R}^m \rightarrow \mathbb{R})$

Torch Core

- The Torch core consists of the following packages:
 - [torch](#): tensors, class factory, serialization, BLAS
 - [nn](#): neural network Modules and Criteria
 - [optim](#): SGD and other optimization functions
 - [gnuplot](#): plotting and data visualization
 - [paths](#): make directories, concatenate file paths, and other filesystem utilities
 - [image](#): save, load, crop, scale, warp, translate image and such
 - [trepl](#): the torch LuaJIT interpreter
 - [cwrap](#): used for wrapping C/CUDA functions in Lua

A snapshot of Torch code

```
In [ ]: net = nn.Sequential()
net.add(nn.SpatialConvolution(1, 6, 5, 5)) -- 1 input image channel, 6 output channels, 5x5 convolution kernel
net.add(nn.SpatialMaxPooling(2,2,2,2))      -- A max-pooling operation that looks at 2x2 windows and finds the max.
net.add(nn.SpatialConvolution(6, 16, 5, 5))
net.add(nn.SpatialMaxPooling(2,2,2,2))
net.add(nn.View(16*5*5))                    -- reshapes from a 3D tensor of 16x5x5 into 1D tensor of 16*5*5
net.add(nn.Linear(16*5*5, 120))              -- fully connected layer (matrix multiplication between input and weights)
net.add(nn.Linear(120, 84))
net.add(nn.Linear(84, 10))                  -- 10 is the number of outputs of the network (in this case, 10 digits)
net.add(nn.LogSoftMax())                    -- converts the output to a log-probability. Useful for classification problems

print('Lenet5\n' .. net.__toString());
```


A quick start with Theano

What is Theano?

- Developed and used since January 2008, created at Universite de Montreal
- Tutorials and examples: <http://deeplearning.net/tutorial/>
- A mathematical symbolic expression compiler
- A Python library for symbolic maths
 - far broader than just Deep Learning
- Tightly integrated with the Python ecosystem
- Fast C/CUDA back-end and transparent GPU acceleration

Theano Recipe

- Recipe for a Theano application:
 - Define symbolic expressions
 - Compile a function that can compile numeric values using those expressions
 - Execute that function on data

Theano Example

$$y = a \times b$$

$$a, b \in \mathbb{R}$$

```
1 import theano
2 from theano import tensor as T
3
4 a = T.scalar()
5 b = T.scalar()
6
7 y = a * b
8
9 multiply = theano.function(inputs=[a, b], outputs=y)
10
11 print multiply(3, 2) #6
12 print multiply(4, 5) #20
13
```

Initialize symbolic variables

Define symbolic expression

Compile a function

Execute on numeric data

Related Projects (ML)

- Keras
- Lasagne
- Pylearn2
- PyMC 3
- sklearn-theano
- theano-rnn
- more...

Built on top of Theano

Typically simplify syntax and
interface for NN training

CNN example in Keras



```
1 from keras.models import Sequential
2 from keras.layers.core import Dense, Dropout, Activation, Flatten
3 from keras.layers.convolutional import Convolution2D, MaxPooling2D
4 from keras.optimizers import SGD
5
6 model = Sequential()
7 model.add(Convolution2D(32, 3, 3, 3, border_mode='full'))
8 model.add(Activation('relu'))
9 model.add(MaxPooling2D(poolsize=(2, 2)))
10
11 model.add(Flatten())
12 model.add(Dense(64*8*8, 256))
13 model.add(Activation('relu'))
14 model.add(Dropout(0.5))
15
16 model.add(Dense(256, 10))
17 model.add(Activation('softmax'))
18
19 sgd = SGD(lr=0.1, decay=1e-6, momentum=0.9, nesterov=True)
20 model.compile(loss='categorical_crossentropy', optimizer=sgd)
21
22 model.fit(X_train, Y_train, batch_size=32, nb_epoch=1)
```

A quick start with TensorFlow

TensorFlow vs. Theano

- Both are deep learning libraries with Python wrapper
- Theano came out first (was inspiration for TensorFlow)
- TensorFlow has better support for distributed systems
- TensorFlow has development funded by Google, while Theano is an academic project
- Everything about TensorFlow is here:
 - <https://www.tensorflow.org>

TensorFlow Essentials

- Four types of objects make TensorFlow unique from other frameworks
 - Session
 - Computational graph
 - Variables
 - Placeholder

TensorFlow Recipe

- Recipe for a TensorFlow application:
 - Define a series of expressions
 - Initialize variables
 - Start a session (launch a graph)
 - Run the graph, feed some data, fetch some values

TensorFlow Session

- “A Session object encapsulates the environment in which Tensor objects are evaluated.” — [TensorFlow Docs](#)

```
In [4]: a = tf.constant(5.0)
        b = tf.constant(6.0)
        c = a*b
```

```
In [5]: print c
```

```
Tensor("mul_2:0", shape=TensorShape([]), dtype=float32)
```

```
In [6]: with tf.Session() as sess:
        print (sess.run(c))
```

```
30.0
```

TensorFlow Computational Graph

- “TensorFlow programs are usually structured into a construction phase, that assembles a graph, and an execution phase that uses a session to execute ops in the graph.” — [TensorFlow Docs](#)

```
In [4]: a = tf.constant(5.0)
        b = tf.constant(6.0)
        c = a*b
```

graph is defined

```
In [5]: print c
```

```
Tensor("mul_2:0", shape=TensorShape([]), dtype=float32)
```

```
In [6]: with tf.Session() as sess:
        print (sess.run(c))
```

graph is launched

```
30.0
```

TensorFlow Variables

- TensorFlow Variables: hold and update **parameters**

```
In [9]: W1 = tf.ones((2,2))  
        W2 = tf.Variable(tf.zeros((2,2)))
```

```
In [10]: init = tf.initialize_all_variables()
```

```
In [11]: with tf.Session() as sess:  
        print sess.run(W1)  
        sess.run(init)  
        print sess.run(W2)
```

```
[[ 1.  1.]  
 [ 1.  1.]  
 [[ 0.  0.]  
 [ 0.  0.]
```

Update Variables

```
In [12]: state = tf.Variable(0)
         new_value = tf.add(state, tf.constant(1))
```

```
In [13]: update = tf.assign(state, new_value)
```

```
In [14]: with tf.Session() as sess:
         sess.run(tf.initialize_all_variables())
         print(sess.run(state))
         for _ in range(5):
             sess.run(update)
             print(sess.run(state))
```

0
1
2
3
4
5

TensorFlow Placeholders

- TensorFlow placeholders: dummy nodes that provide entry points for data to computational graph

```
In [15]: input1 = tf.placeholder(tf.float32)
         input2 = tf.placeholder(tf.float32)
```

```
In [16]: output = tf.mul(input1, input2)
```

```
In [20]: with tf.Session() as sess:
         print(sess.run(output, feed_dict={input1:7., input2:2.}))
```

14.0

Thank you!

Feedbacks and Questions: rosanne.liu@northwestern.edu